

Participation Costs, Social Capital, and the Limits of Mobility Policy

A Model of Status Competition under Income Inequality

Anurag Srivastava

Riskcare Ltd.

Email: anurag.srivastava@riskcare.com ORCID: 0000-0002-6477-4430

Corresponding author

May 21, 2026

Abstract

Widening access to elite competitions can reinforce, rather than erode, class barriers. We show this as the generic equilibrium outcome once participation is treated as an endogenous, income-constrained decision. A high-income incumbent competes against low-income challengers in tournaments rewarding both innate ability and acquired social capital, where the latter requires a threshold expenditure. Exactly two stable Nash equilibria arise – universal and elite-only participation – and a threshold characterises which prevails. Wider income gaps simultaneously strengthen the incumbent’s investment incentive and price out challengers, reconciling rising inequality with rising status expenditure. The central finding is a hard mobility ceiling governed by the weight placed on raw talent relative to social capital: no reduction in participation costs, however large, can breach it without a simultaneous change in what the competition rewards.

Keywords: status competition; income inequality; social capital; participation threshold; mobility ceiling; public policy

JEL classification: D31, J62, Z13

1 Introduction

Whether income inequality intensifies or dampens competition for social position is an old question with a surprisingly durable empirical answer. Post-war scholars — Hirsch (1977), Galbraith (1958), and Frank (1985) — predicted that rising living standards and narrowing income gaps would intensify positional rivalry, as more participants could afford to compete for a fixed stock of social rank. Contemporary evidence consistently finds the opposite: status-related expenditure rises with inequality rather than falling with it, with education expenditure in China [Jin et al., 2011] or conspicuous

goods across multiple countries [Chai et al., 2019, Jaikumar and Sarin, 2015] providing systematic documentation.

This paper concerns itself with the specific form of status expenditure that has direct consequences for economic outcomes: investment in educational credentials and institutional affiliation. Credentials are positional goods in the sense of Hirsch [1977]: their competitive value derives from the scarcity of the positions they unlock, not solely from the human capital they embody. In Bourdieu’s framework [Bourdieu and Passeron, 1977, Bourdieu, 1986], credentials function as institutionalised cultural capital: they confer competitive advantage in selection processes through the social recognition they command rather than solely through the productivity they embody. Rising inequality raises both the stakes of winning a credentialled position and the structural advantage of those who can afford the threshold investment required to compete — which is why education expenditure and inequality may move together in the data. The present model treats education as such a credential rather than as a productive asset that accumulates individual capacity across jobs. This focus reflects the observed reality that access to economic goods — elite occupations, professional status, income at the top of the distribution — is often allocated through positional competition rather than anonymous market clearing [Rivera, 2015, Laurison and Friedman, 2016, Bourdieu, 1986]: credentialism is the mechanism, not an aberration from it.

The source of the discrepancy, we argue, lies in a common simplification: existing models of status competition treat the participation decision as a continuous intensity choice — agents decide how much to invest in status-enhancing social capital, with income affecting the level of investment but not whether entry is feasible at all. A threshold cost structure that makes participation a discrete binary decision — affordable or not, depending on income — allows us to inspect the effects of income inequality on binary outcomes of success and shows that wider income gaps simultaneously strengthen the high-income participant’s investment incentive and price out low-income challengers. The positive association between inequality and status expenditure follows as the generic equilibrium outcome, not a special case.

The model’s specific architecture — the separation of talent μr_i from social capital $(1 - \mu)s_i$ — the talent-weighted and social-capital-weighted components respectively — of competitive rank — is motivated by several empirical patterns in the elite mobility literature that the standard Becker-Tomes human capital framework doesn’t generate without auxiliary assumptions about non-productive forms of capital. Lo Bello and Morchio [2022] find that dynastic workers in the UK who follow their parent’s occupation find jobs faster but earn *lower* wages than comparably qualified non-dynastic workers; human capital theory predicts wage parity or a premium, not a penalty. Laurison and Friedman [2016] show that the class-origin pay gap in British elite occupations varies sharply by occupation type: it is close to zero in science, academia, and the built environment but approximately 20% in finance, law, and accountancy — the opposite of a human capital view since the gap should be largest where inherited knowledge is most valuable. The data are directly consistent with the model’s prediction that the class gap is increasing in $(1 - \mu)$ (social capital weight): finance and law are low- μ (talent weight) occupations where networks and cultural fit dominate, while science and engineering are high- μ (talent weight) occupations where technical output is directly measurable. Mocetti et al. [2022] exploit two waves of liberalisation of professional services in Italy and

find that deregulation reduces intergenerational occupational persistence, with the effect concentrated among less able dynastic workers (proxied by delayed school qualification) — consistent with the model’s prediction that social capital weight $(1 - \mu)$ determines how far raw talent shortfalls can be compensated by positional advantages, so that low- r_i incumbents are the first to exit when regulatory protection shrinks.¹ These residuals share a common structure: each identifies a component of elite access that facilitates entry without raising productivity, and whose weight varies across occupations. That component is the social capital weight $(1 - \mu)$ in the present model.

Against this backdrop, the paper develops the following model. The model considers a status competition in which one high-income participant competes against $Q \geq 1$ low-income participants through a probabilistic tournament combining innate ability and acquired social capital. Social capital — the skills, credentials, and network affiliations that institutional engagement provides — can be obtained only by paying a threshold cost $\bar{v}$, making it accessible to the high-income participant by default and contingent on budget feasibility for low-income participants. The probability of advancement depends on both components, with social capital weighted more heavily than raw talent, following the empirical literature on elite hiring [Rivera, 2015, Bills, 2003]. Status enters utility indirectly through its effect on future income, following Postlewaite [1998]; no separate taste for rank is required.

The model yields four results with direct policy content. First, under diminishing wealth utility only two stable Nash equilibria exist: universal participation (both invest) and elite-only participation (only the high-income participant invests). The transition between them is governed by a participation threshold $\bar{v}^*$ that decreases in both competition size Q and the social capital weight $(1 - \mu)$. Wider inequality sustains high-income investment even as it prices out low-income challengers, which is why status expenditure and inequality rise together.

Second, collective low-income mobility converges to a hard ceiling of $\mu/[2(1 - \mu)]$ in the empirically relevant regime. This ceiling is bounded strictly below one-half for all admissible μ , and falls as social capital is weighted more heavily. It represents a structural limit that no participation-boosting policy can overcome without also raising μ (talent weight). Reducing participation costs alone expands the range of inequality levels consistent with universal participation, but the ceiling itself is determined by the design of the competition, not the number of participants.

Third, these two instruments are complements in a precise sense: the participation threshold $\bar{v}^*$ is increasing in A , a composite parameter embedding both Q and μ (talent weight), and the ceiling $\mu/[2(1 - \mu)]$ is increasing in μ (talent weight). Policies that lower $\bar{v}$ (credential subsidies, open-access qualifications, reduced network barriers) shift the equilibrium toward universal participation; policies that raise μ , the talent weight (standardised assessment, transparent merit criteria), raise the ceiling at which collective mobility converges. Neither alone is sufficient.

Fourth, a third equilibrium — in which only low-income participants invest while the

¹A related finding concerns UK university admissions, where personal statements — a threshold social capital investment requiring coaching and institutional support — have negligible predictive validity for university performance [Jones, 2013], and students from less advantaged backgrounds out-perform peers from more advantaged backgrounds with the same entry grades. This is consistent with the model’s architecture but pertains specifically to education admissions rather than labour market competition.

high-income participant free-rides on raw talent — requires reference-dependent preferences inverting the standard concavity ordering, and is viable only in small competitions ($Q = 1$ or 2). This provides a formal account of the over-investment in elite credentials observed among aspiring candidates from lower-income backgrounds in narrow selection processes [Rivera, 2015]: the mechanism is both preference-dependent and competition-size-specific, and cannot survive scaling to large standardised assessments.

These results position the paper relative to its quantitative predecessors as follows. The two closest — each reaching conclusions that diverge from the present paper — fix the participant pool. Frank [1984] shows that wider wage dispersion reduces competition intensity — workers accept wage penalties to occupy the top of the local rank distribution, so narrowing inequality intensifies rivalry. Cole et al. [1992] derive status concern from the marriage market and show directly that agents from a more compact income distribution save more, again implying that narrowing inequality raises status expenditure. In both cases income distribution shapes the intensity of competition within the participant pool, but participation itself is not subject to a threshold budget constraint — all agents compete, and the question is how intensely. Both models therefore predict that *narrowing* inequality raises status expenditure, the opposite of the empirical finding. The present paper introduces a threshold participation cost that makes entry a discrete binary decision, and shows that this single change reverses the sign: *widening* inequality raises status expenditure by sustaining high-income investment while pricing out low-income challengers.

The most directly comparable model is that of Gallice and Grillo [2019], who consider a model in which educational investment entails productivity gains, signaling power, and social status simultaneously, with status depending on relative achievement in one of three dimensions: innate skills, level of schooling, and income. They characterise conditions under which social concerns increase or decrease educational and income inequality. The present paper asks a different question: not how much agents invest conditional on universal participation, but whether participation itself is feasible across income levels. The current paper generates a structural bound and allows hierarchy-related conclusions for mobility using the ceiling $\mu/[2(1 - \mu)]$ and the binary equilibrium classification (EP–FP) which arise specifically from the threshold cost structure and income asymmetry — claims that don’t follow from the Gallice and Grillo model — where all agents participate with varying investment intensity. Whether inequality amplifies or dampens status investment conditional on universal participation is a distinct question from whether inequality sustains or destroys participation itself.

The threshold cost structure used in the paper is motivated by broader literature on social capital and credential investment [Jencks, 1979, Heckman and Mosso, 2014]. Participation costs in elite competitions — tuition fees, membership dues, network cultivation — are threshold rather than continuous expenditures: they must be paid in full to unlock the social capital they confer. This threshold structure is central to the model’s ability to generate binary equilibrium outcomes and a closed-form participation condition. The sociological literature [Wegener, 1992, Mijs, 2021, Mijs and Savage, 2020] motivates the empirical puzzle and the inclusive-participation (IP) over-investment result — following from standard expected-utility maximisation with a threshold entry cost.

The paper proceeds as follows. Section 2 develops the model, deriving winning

probabilities for all four investment cases under general Q , characterising the Nash equilibria, deriving the participation threshold in closed form, and establishing the mobility ceilings. Section 3 draws out the policy implications, mapping the model’s parameters onto observable competition structures. Section 4 concludes.

2 Model

2.1 Basic Setup

We consider a status competition in which a single high-income participant competes against $Q \geq 1$ low-income participants ($Q + 1$ total). The high-income participant occupies the high-income position y_2 and seeks to maintain it; each low-income participant receives income $y_1 < y_2$ and seeks promotion. Two characteristics determine performance in the competition. First, *raw talent* $r_i \sim \text{Uniform}[0, 1]$, which cannot be affected by expenditure. Second, *social capital* $s_i \sim \text{Uniform}[0, 1]$, which can be enhanced only by paying a threshold cost $\bar{v}$. Both r_i and s_i are independently distributed and unknown to participant i before the competition. The participant with the highest composite score $\mu r_i + (1 - \mu)s_i$ wins the high-income position, where $\mu \in (0, 1)$ weights raw talent relative to social capital.

Since all Q low-income participants share the same income y_1 and all face the same investment decision, we analyze the participation problem from the perspective of a representative low-income participant. The high-income participant’s decision is analogous. A participant pays $\bar{v}$ to access social capital (setting $s_i > 0$) or abstains and saves instead (setting $s_i = 0$, relying on raw talent alone). Crucially, no participant knows r_i or s_i prior to deciding whether to pay $\bar{v}$, so investment decisions depend only on the income class of the participant.

Status enters utility only indirectly: the winner obtains disposable income y_2 in the following period, the loser retains y_1 . This treatment, following Postlewaite [1998], preserves a standard wealth-utility framework while endogenising the social dimension of the competition.

The model is designed for a specific class of competitions: those in which a fixed number of high-income positions is contested by a variable pool of challengers, and in which the winner’s gain is the loser’s foregone advancement. This zero-sum structure is not a simplifying approximation but an accurate description of positional competitions for elite credentials — partnership tracks at professional firms, senior academic appointments, competitive university admissions — where the number of positions is institutionally fixed and occupying one precludes others from doing so. The model does not speak to general labour markets in which education raises productivity across all jobs, nor to settings in which a losing investor retains a human capital gain from the investment. That class of questions — where social capital raises y_1 as well as the probability of winning y_2 — is a natural extension in which the threshold cost generates a productivity externality for non-winners; we leave it to future work. The present model isolates the participation and mobility effects of threshold costs in positional competitions, which is the mechanism most directly relevant to the empirical puzzle that motivates it.

Focusing on conditions where the contest advantage from social capital investment is substantial, we assume throughout that social capital is weighted more heavily than raw talent:

Assumption 1. $\mu < \frac{1}{2}$

Remark on distributional assumptions. The uniform distributions for r_i and s_i are adopted for tractability: they yield the closed-form trapezoidal CDF $F_X(x)$ whose piecewise structure drives the winning probability derivations in Appendix A. The key qualitative results do not depend on uniformity. The fundamental ordering $p_1 < 1/(Q+1) < p'_1$ — which drives the Nash equilibrium analysis — holds for any pair of distributions that are symmetric on $[0, 1]$ and have the same marginal distribution for r_i and s_i , since the ordering follows from the structural disadvantage of competing without social capital against an opponent who has it, not from the specific shape of the distributions. The closed-form expressions for p_1 , the participation threshold $\bar{v}^*$, and the mobility ceiling $\mu/[2(1-\mu)]$ are specific to the uniform case, but the existence of the binary equilibrium structure (exclusive participation, EP, and full participation, FP, as the only generically stable equilibria under diminishing utility), and all comparative statics on Q , μ , $\bar{v}$, and $y_2 - y_1$ extend to the broader distributional class. The uniform specification thus functions as a parametric representation of a more general result rather than a special case on which the argument depends.

The distribution of the composite score $x_i \equiv \mu r_i + (1-\mu)s_i$ under Assumption 1 is trapezoidal (derived in Appendix A), with CDF $F_X(x)$ that is piecewise quadratic on $[0, \mu]$, linear on $[\mu, 1-\mu]$, and piecewise quadratic on $[1-\mu, 1]$.

The four investment cases are as follows.

- **Exclusive Participation (EP)** — High-income invests ($v_2 = \bar{v}$), low-income does not ($v_1 = 0$).
- **Inclusive Participation (IP)** — Low-income invest ($v_1 = \bar{v}$), high-income does not ($v_2 = 0$).
- **Full Withdrawal (FW)** — Neither invests ($v_1 = v_2 = 0$).
- **Full Participation (FP)** — Both invest ($v_1 = v_2 = \bar{v}$).

2.2 Equilibrium Conditions

2.2.1 Winning Probabilities

The derivations are in Appendices A and B. The per-low-income-participant winning probabilities are:

$$\begin{aligned}
(1) \quad p_1 &= \frac{\mu}{2(1-\mu)(Q+2)} & (\text{EP}) \\
(2) \quad p'_1 &= \frac{6-7\mu}{6(1-\mu)} & (\text{IP}, Q=1) \\
(3) \quad p_1 &= \frac{1}{Q+1} & (\text{FW and FP})
\end{aligned}$$

In EP, the low-income participant competes on raw talent alone against the high-income participant's full composite score — a structural disadvantage that drives p_1 strictly below the symmetric baseline $1/(Q+1)$. In IP the roles are reversed: the investing low-income participant has the full score advantage, pushing p'_1 strictly above the baseline. For FW and FP, symmetry gives $1/(Q+1)$ directly. For general $Q \geq 2$, p'_1 is computable as a rational function of μ of degree Q (Appendix B); all equilibrium results for IP follow from the ordering result alone, without requiring a closed form for p'_1 .

The high-income participant's probability follows from $p_2 = 1 - Qp_1$ in each case.

The fundamental ordering that drives the Nash equilibrium analysis follows directly from Assumption 1:

$$(4) \quad p_1 < \frac{1}{Q+1} < p'_1 \quad \text{for all } Q \geq 1, \mu < \frac{1}{2}$$

Figure 1 illustrates these probabilities. Panel (a) shows that p_1 (EP) lies strictly below the baseline $1/(Q+1)$ and decays rapidly with Q . Panel (b) shows p_2 (EP, high-income participant) lying well above the baseline. Panel (c) shows p'_1 (IP, $Q=1$) strictly above the baseline throughout. Panel (d) is discussed in Section 2.2.5.




Figure 1: Winning probabilities and mobility ceiling. One high-income participant, Q low-income participants. Dotted horizontal lines in panels (a)–(b) show baselines $1/(Q+1)$ for each Q .

2.2.2 Intertemporal Choice and Utility Payoffs

Participants optimise an intertemporally additive utility $u(\cdot)$ over T periods. In period 1 each participant decides whether to pay $\bar{v}$; from period 2 onward she receives the income corresponding to her position. Define the T -period discount sum $D_T = \sum_{t=0}^{T-2} \delta^{-t} = (1 - \delta^{-(T-1)})/(1 - \delta^{-1})$, where $\delta > 1$ is the discount factor. The case $T = 2$ gives $D_T = 1$ and recovers the standard two-period model. The expected lifetime utilities are:

$$\text{Low-income participant: } u(y_1 - v_1) + \frac{D_T}{\delta} [p_1^c u(y_2) + (1 - p_1^c) u(y_1)]$$

$$\text{High-income participant: } u(y_2 - v_2) + \frac{D_T}{\delta} [p_2^c u(y_2) + (1 - p_2^c) u(y_1)]$$

where p_i^c denotes the winning probability in case $c \in \{\text{EP}, \text{IP}, \text{FW}, \text{FP}\}$. Define the utility costs of participation:

$$U_1 \equiv u(y_1) - u(y_1 - \bar{v}), \quad U_2 \equiv u(y_2) - u(y_2 - \bar{v})$$

Under diminishing marginal utility with $y_1 < y_2$: $U_2 < U_1$. Table 1 summarises the four cases as a 2×2 game; full utility payoffs are in Table 4 (Appendix D).

Table 1: Investment strategies

	<i>Low-income withdraws</i>	<i>Low-income invests</i>
<i>High-income withdraws</i>	FW	IP
<i>High-income invests</i>	EP	FP

The full utility payoffs for each case are reported in Table 4 in Appendix D; the equilibrium conditions below are derived from those payoffs.

2.2.3 Nash Equilibrium Conditions

For ease of notation define the probability-gap terms:

$$(5) \quad \begin{aligned} p_1 &\equiv \frac{\mu}{2(1-\mu)(Q+2)}, \quad p_2 \equiv 1 - Qp_1 \\ p'_1 &\equiv \frac{6-7\mu}{6(1-\mu)} \quad (Q=1) \end{aligned}$$

The ordering $p_1 < 1/(Q+1) < p_2$ and $p'_1 > 1/(Q+1)$ holds for all $\mu < 1/2$ (Equation 4). Note that $p'_2 = 1 - Qp'_1$ satisfies $p'_2 < 1/(Q+1)$ by the same ordering, but p'_2 does not appear in the equilibrium conditions below: by the algebraic identity $1/(Q+1) - p'_2 = Q(p'_1 - 1/(Q+1))$, it is replaced throughout by $Q \cdot \Phi(Q)$, which is valid for all $Q \geq 1$. We define $\Phi(Q)$ now, before analysing the cases, since it appears in both EP and IP:

The *probability surplus* $\Phi(Q)$ is the gain in winning probability a low-income participant obtains from being the sole investor in social capital relative to the symmetric no-investment baseline, scaled by the utility gap and discount factor:

$$(6) \quad \Phi(Q) \equiv \frac{D_T}{\delta} \left(p'_1 - \frac{1}{Q+1} \right) (u(y_2) - u(y_1))$$

Since $p'_1 > 1/(Q+1)$ for all $\mu < 1/2$ (Equation 4), $\Phi(Q) > 0$ always. For $Q = 1$, substituting $p'_1 = (6-7\mu)/[6(1-\mu)]$ yields:

$$(7) \quad \Phi(1) = \frac{D_T}{\delta} \cdot \frac{3-4\mu}{6(1-\mu)} \cdot (u(y_2) - u(y_1))$$

where $(3-4\mu) > 0$ under Assumption 1, so $\Phi(1)$ is strictly positive and decreasing in μ , ranging from $\Phi(1) \rightarrow 1/2$ at $\mu \rightarrow 0$ to $\Phi(1) \rightarrow 1/3$ at $\mu \rightarrow 1/2$. By Proposition 2 below, $\Phi(Q) \sim 1/[Q(Q+1)]$ asymptotically, so the surplus narrows rapidly with competition size.

Exclusive Participation (EP) (high-income invests, low-income withdraws) is a Nash equilibrium if and only if:²

$$(8) \quad \frac{D_T}{\delta} \left(\frac{1}{Q+1} - p_1 \right) (u(y_2) - u(y_1)) < U_1$$

$$(9) \quad U_2 < \frac{D_T}{\delta} \left(p_2 - \frac{1}{Q+1} \right) (u(y_2) - u(y_1))$$

Both conditions hold simultaneously under diminishing utility since $U_2 < U_1$ and $p_2 - 1/(Q+1) \geq 1/(Q+1) - p_1$. A higher $\bar{v}$ favors EP so long as the high-income participant retains an incentive to invest; a lower $\bar{v}$ can destabilize it if low-income participants are induced to enter. Wider income differences ($y_2 - y_1$ larger) reinforce the high-income participant's incentive to invest while simultaneously raising the stakes enough to attract low-income investment, making the first condition harder to satisfy.

Inclusive Participation (IP) (low-income invest, high-income withdraws). We now state this case as a formal proposition, since it requires conditions beyond standard diminishing utility and its characterisation carries independent theoretical content. The governing object is $\Phi(Q)$, defined above. The probability surplus narrows at rate Q^{-2} : at $Q = 1$ the IP window is wide ($\Phi(1) \approx 0.46$ at $\mu = 0.2$), but by $Q = 3$ it has collapsed to $\Phi(3) \approx 0.08$ — an 82% reduction — and by $Q = 5$ to $\Phi(5) \approx 0.03$ (Table 3). The qualitative conclusions of Proposition 1 — that IP requires reference-dependent preferences and that the feasibility window $[U_1, U_2/Q]$ narrows as μ increases — hold for general $Q \geq 1$ by the ordering result in Equation 4.

Proposition 1 (IP Nash Equilibrium). *IP (low-income invest, high-income withdraws) is a Nash equilibrium if and only if:*

$$(10) \quad U_1 < \Phi(Q) \leq U_2/Q$$

A necessary condition is therefore $Q \cdot U_1 < U_2$. Under standard diminishing utility with $y_1 < y_2$, $U_2 < U_1$, so $U_2/Q < U_1$ for all $Q \geq 1$ and Condition (10) cannot be satisfied. IP requires reference-dependent preferences under which $U_1 < U_2$ — the marginal utility cost of $\bar{v}$ is higher for the low-income participant than for the high-income participant, reversing the standard concavity ordering. For $Q = 1$, the condition reduces to:

$$(11) \quad U_1 < \Phi(1) \leq U_2$$

which is feasible if and only if $U_1 < U_2$. For general Q , the necessary condition $Q \cdot U_1 < U_2$ tightens as Q grows; by Proposition 2, $\Phi(Q) \sim 1/[Q(Q+1)]$, so the condition fails for any fixed $U_1 > 0$ beyond $Q^* \approx 1/\sqrt{U_1}$.

Proof. Condition (10) is derived directly from the utility payoffs in Table 4 (Appendix D). The low-income participant prefers to invest given the high-income participant withdraws if and only if $U_1 < \Phi(Q)$. The high-income participant prefers to

²Derived from the payoff table (Appendix D): the low-income participant prefers not to invest, and the high-income participant prefers to invest, given the other's strategy.

withdraw given the low-income participants invest if and only if the gain from deviating to investment, $\frac{D_T}{\delta} (1/(Q+1) - p'_2)(u(y_2) - u(y_1)) = Q \cdot \Phi(Q)$ (using $1/(Q+1) - p'_2 = Q(p'_1 - 1/(Q+1))$), valid for all $Q \geq 1$, is less than the cost U_2 , i.e. $U_2 \geq Q \cdot \Phi(Q)$, equivalently $\Phi(Q) \leq U_2/Q$. Both conditions together require $U_1 < U_2/Q$. Under standard concavity $U_2 < U_1$, so $U_2/Q < U_1$ for all $Q \geq 1$ and the condition fails. $\square \quad \square$

Several features of this proposition deserve comment. First, $\Phi(1)$ is strictly decreasing in μ (talent weight): as social capital is weighted more heavily, the probability surplus from being the sole investor narrows, and the IP window must be correspondingly tighter to sustain the equilibrium. At $\mu \rightarrow 0$ (pure raw-talent competition), $\Phi(1) \rightarrow 1/2$, giving the widest possible window; at $\mu \rightarrow 1/2$ (Assumption 1 binding), $\Phi(1) \rightarrow 1/3$. Second, the equilibrium remains fragile even when it exists: higher $\bar{v}$, greater impatience (δ rising), or wider income differences ($y_2 - y_1$ larger) can each collapse it into FW or FP respectively. Third, $\Phi(Q)$ narrows at rate Q^{-2} for general Q (Proposition 2), making IP quantitatively relevant only for small competitions ($Q = 1$ or 2). This is consistent with the interpretation that IP — over-investment by aspiring low-income candidates in socially embedded competitions — is a phenomenon of elite selection processes with narrow fields, not of large-scale standardised assessments.

Full Withdrawal (FW) (neither invests) is a Nash equilibrium if:³

$$\begin{aligned} U_1 &> \frac{D_T}{\delta} \left(p'_1 - \frac{1}{Q+1} \right) (u(y_2) - u(y_1)) \\ U_2 &> \frac{D_T}{\delta} \left(p_2 - \frac{1}{Q+1} \right) (u(y_2) - u(y_1)) \end{aligned}$$

Since $p'_1 - 1/(Q+1) \leq p_2 - 1/(Q+1)$ and $U_2 < U_1$, the binding constraint is the second inequality. FW prevails when participation costs are prohibitively high relative to the income gain. If there is any net utility from competition, FP is the likely outcome.

Full Participation (FP) (both invest) is a Nash equilibrium if:⁴

$$\begin{aligned} (12) \quad U_1 &< \frac{D_T}{\delta} \left(\frac{1}{Q+1} - p_1 \right) (u(y_2) - u(y_1)) \\ (13) \quad U_2 &< Q \cdot \Phi(Q) \end{aligned}$$

With $U_2 < U_1$, this equilibrium rests on the low-income participant's incentive: so long as the low-income participant finds the competition worthwhile, the high-income participant also participates. Note that Equations 8 and 12 are exact complements — FP obtains precisely where EP fails for the low-income participant. The high-income participant's condition $U_2 < Q \cdot \Phi(Q)$ uses the algebraic identity $1/(Q+1) - p'_2 = Q(p'_1 - 1/(Q+1))$, which holds for all $Q \geq 1$ and connects FP directly to the probability

³Both participants prefer withdrawal given the other withdraws.

⁴Both participants prefer investment given the other invests.

surplus $\Phi(Q)$ from Proposition 1: the high-income participant's boundary between IP and FP is exactly $Q \cdot \Phi(Q)$.

These conditions confirm that under diminishing utility only EP and FP are generically stable. For very low $\bar{v}$, both participants invest (FP); for very high $\bar{v}$, only the high-income participant invests (EP); for sufficiently extreme $\bar{v}$, neither invests (FW). IP requires a non-standard utility structure and remains fragile throughout.

2.2.4 Participation Threshold

The boundary between FP and EP is defined by the critical cost $\bar{v}^*$ at which the low-income participant is indifferent between investing and withdrawing, given the high-income participant invests. Setting Equation 12 as an equality:

$$(14) \quad u(y_1) - u(y_1 - \bar{v}^*) = \frac{D_T}{\delta} \left(\frac{1}{Q+1} - p_1 \right) (u(y_2) - u(y_1))$$

Under log utility $u(y) = \log(y)$, this yields $\bar{v}^* = y_1(1 - (y_1/y_2)^A)$, showing that the threshold depends on the income ratio y_1/y_2 rather than the income level alone; the full derivation is in Appendix C.

Here, the *participation window parameter* A is defined as:

$$(15) \quad A \equiv \frac{D_T}{\delta} \left(\frac{1}{Q+1} - \frac{\mu}{2(1-\mu)(Q+2)} \right) > 0$$

$A > 0$ for all $\mu < 1/2$ and $Q \geq 1$ since $1/(Q+1) > \mu/[2(1-\mu)(Q+2)]$ under Assumption 1. The threshold $\bar{v}^*$ is increasing in A , so any comparative static that reduces A tightens the participation window.

Competition size (Q increasing, more low-income challengers):

$$(16) \quad \frac{\partial A}{\partial Q} = \frac{-\mu(Q+1)^2 + 2(1-\mu)(Q+2)^2}{2(Q+1)^2(Q+2)^2(1-\mu)} \cdot \frac{1}{\delta}$$

Since $\mu < 1$, the denominator $2(Q+1)^2(Q+2)^2(\mu-1) < 0$, while the numerator $N = 2(1-\mu)(Q+2)^2 - \mu(Q+1)^2$ satisfies $N > (Q+2)^2 - \frac{1}{2}(Q+1)^2 = \frac{1}{2}Q^2 + 3Q + \frac{7}{2} > 0$ for all $Q \geq 0$ under $\mu < \frac{1}{2}$. Hence $\partial A / \partial Q < 0$. As the pool of low-income challengers grows, A decreases and $\bar{v}^*$ falls. Equivalently, a larger Q makes EP increasingly likely to prevail over FP. As $Q \rightarrow \infty$, $A \rightarrow 0$ and the participation window closes entirely.

Competition size and real-world competition structures. The comparative statics on Q and μ (talent weight) generate a sharp distinction between small competitions (low Q , low μ) — partner-track interviews, senior academic appointments, narrow graduate schemes — and large competitions (high Q , higher μ) — civil service examinations, mass graduate recruitment, standardised professional qualifications. The empirical interpretation of these comparative statics and their implications for the design of each competition type are developed in Section 3.

2.2.5 Collective Low-Income Probability and the Mobility Ceiling

The analysis so far has focused on the per-low-income-participant winning probability p_1 . A distinct and complementary object is the *collective low-income probability* $Q \cdot p_1$, which measures the probability that any low-income participant displaces the high-income participant — a statement about class-level mobility rather than individual opportunity. Two ceilings emerge depending on which equilibrium prevails.

EP ceiling. Under EP, the collective low-income probability in closed form is:

$$(17) \quad Q \cdot p_1 = \frac{Q\mu}{2(1-\mu)(Q+2)}$$

As $Q \rightarrow \infty$:

$$(18) \quad \lim_{Q \rightarrow \infty} Q \cdot p_1 = \frac{\mu}{2(1-\mu)}$$

This ceiling is bounded strictly below $\frac{1}{2}$ for all $\mu < \frac{1}{2}$ and is increasing in μ (talent weight).

FP ceiling. Under FP (both invest), all $Q+1$ participants draw from the same distribution F_X , so $p_1 = 1/(Q+1)$ and:

$$(19) \quad \lim_{Q \rightarrow \infty} Q \cdot \frac{1}{Q+1} = 1$$

Under universal participation, collective low-income mobility approaches certainty as the field grows.

Why the EP ceiling is the economically relevant one. The two ceilings appear to give contradictory messages about the social limits to growth. The resolution lies in the comparative statics: by Equation 16, $\partial A / \partial Q < 0$, so a larger pool of low-income challengers *tightens* the participation window and makes EP increasingly likely to prevail over FP. The EP ceiling is therefore precisely the one that binds as Q grows: the equilibrium converges toward elite-only participation at exactly the same rate as collective mobility approaches its hard limit, and simultaneously the high-income participant's FP incentive condition $U_2 < Q \cdot \Phi(Q)$ weakens as $Q \cdot \Phi(Q) \rightarrow 0$ (Corollary 1). The FP ceiling ($\rightarrow 1$) is relevant only when Q is small and costs are low enough to sustain universal participation — the regime where mobility concerns are least acute. It is the EP ceiling $\mu/[2(1-\mu)]$ that constitutes the binding social limit to growth in the sense of Hirsch [1977].

Simultaneously, the *individual* low-income participant's probability $p_1 \rightarrow 0$ as $Q \rightarrow \infty$: each low-income participant's chances vanish even as collective mobility converges to the ceiling. This divergence between individual and collective mobility — the many trying, the class ceiling binding — is the formal content of the social limits argument.

3 Policy Implications

The model's equilibrium structure carries direct implications for the design of education and credential policy. Three themes emerge: the joint configuration of participation costs and competition design that governs which equilibrium prevails; the mobility ceiling that defines the hard limit on class-level advancement regardless of who participates; and the sharp difference in policy leverage between small and large competitions. We address each in turn.

3.1 Participation Costs and Income Differences as Joint Instruments

The transition between the universal participation equilibrium (FP) and the elite-only equilibrium (EP) is governed jointly by $\bar{v}$ and the income gap $y_2 - y_1$. Two regimes apply.

When $y_1 < \bar{v} < y_2$, the budget constraint is binding: participation is ruled out regardless of expected returns, and the elite-only equilibrium is the only feasible outcome. Only a reduction in $\bar{v}$ large enough to restore $y_1 > \bar{v}$ shifts the equilibrium; marginal cost reductions that leave $y_1 < \bar{v}$ do not change the equilibrium and merely lower the income gap at which the constraint becomes binding in future.

When $\bar{v} < y_1 < y_2$, the constraint is non-binding and the participation threshold $\bar{v}^*$ governs the transition. Under log utility:

$$(20) \quad \bar{v}^* = y_1 \left(1 - \left(\frac{y_1}{y_2} \right)^A \right), \quad A \equiv \frac{D_T}{\delta} \left(\frac{1}{Q+1} - \frac{\mu}{2(1-\mu)(Q+2)} \right)$$

Three policy-relevant comparative statics follow directly from Equation (20).

First, $\bar{v}^*$ depends on the income *ratio* y_1/y_2 rather than the income level alone. A uniform upward shift in both incomes that leaves the ratio constant has no effect on the participation threshold. Policies that raise low-income earnings without compressing the income ratio — for example, minimum wage increases in an economy with proportionally rising top incomes — do not expand the participation window. Policies that directly compress y_2/y_1 , or that lower $\bar{v}$ relative to y_1 , are required.

Second, $\bar{v}^*$ is decreasing in Q (since $\partial A / \partial Q < 0$): a larger pool of low-income challengers reduces each individual's expected return from entry and shrinks the range of costs consistent with universal participation. As competition for a fixed number of elite positions becomes more crowded, the rational response of each low-income participant is more likely to be withdrawal, regardless of the income gap. This comparative static is central to understanding why expanding access to elite competitions — increasing the pool of low-income applicants — can paradoxically reinforce elite-only equilibria unless it is accompanied by a simultaneous reduction in $\bar{v}$.

Third, $\bar{v}^*$ is decreasing in μ (talent weight): competitions that weight social capital more heavily reduce the expected return to low-income entry, shrinking the participation window. This formalises the intuition behind credentialist barriers: the more decisive institutional affiliation and credentials are relative to raw talent, the narrower the set of cost and income configurations that sustain universal participation.

3.2 The Mobility Ceiling and Its Policy Implications

Even when the universal participation equilibrium is sustained, collective low-income mobility faces a structural ceiling that participation policy cannot remove. Under the elite-only equilibrium, Exclusive Participation (EP), collective low-income mobility — the probability that any low-income participant displaces the high-income incumbent — converges as $Q \rightarrow \infty$ to:

$$(21) \quad \lim_{Q \rightarrow \infty} Q \cdot p_1 = \frac{\mu}{2(1 - \mu)}$$

This ceiling is bounded strictly below $\frac{1}{2}$ for all admissible $\mu < \frac{1}{2}$, and is increasing in the talent weight μ . Table 2 reports the ceiling for representative values of μ , alongside the corresponding minimum income ratio y_1/y_2 at which the budget constraint becomes non-binding for a standardised threshold cost.

Table 2: Mobility ceiling and participation window by social capital weight

μ	Ceiling $\mu/[2(1 - \mu)]$	Ceiling (%)	Interpretation
0.10	0.056	5.6%	Social capital near-decisive
0.20	0.125	12.5%	High credential weight
0.30	0.214	21.4%	Moderate credential weight
0.40	0.333	33.3%	Talent and credentials balanced
0.49	0.490	49.0%	Near-maximum under Assumption 1

Ceiling = $\mu/[2(1 - \mu)]$ from Equation (21). Values report the long-run probability that the high-income position is held by a low-income-origin participant, as $Q \rightarrow \infty$ under Exclusive Participation (EP).

The empirical relevance of these figures is immediate. Intergenerational income mobility studies for the United Kingdom — including Blanden et al. [2005] — find that children born into the bottom income quartile have approximately a 16% probability of reaching the top income quartile as adults. These figures are consistent with the model’s ceiling under $\mu \in [0.20, 0.30]$, a range that corresponds to competitions in which institutional credentials and network affiliations carry substantial but not overwhelming weight relative to raw ability. The model does not claim to calibrate mobility precisely; rather, it identifies the structural parameter μ (talent weight) that would be necessary to account for observed mobility rates, and shows that this parameter is directly amenable to policy intervention through competition design.

The ceiling $\mu/[2(1 - \mu)]$ is the binding constraint in the empirically relevant regime. Under universal participation, or Full Participation (FP), collective mobility converges to 1 as $Q \rightarrow \infty$ — but this ceiling is only operative when the field is small and participation costs are low enough to sustain universal entry, precisely the conditions least likely to hold in large competitions. It is the Exclusive Participation (EP) ceiling that constitutes the operative social limit to growth as Hirsch (1977) originally argued.

Two policy implications follow directly. The ceiling is determined by μ (talent weight) alone: reducing $\bar{v}$ or expanding the pool of challengers affects which equilibrium

prevails but cannot raise the ceiling itself. Raising μ (talent weight) through standardised, competency-based assessment is therefore a necessary complement to cost reduction. Cost reductions that shift the equilibrium from Exclusive Participation (EP) to Full Participation (FP) produce a discrete jump in the operative ceiling from $\mu/[2(1-\mu)]$ to 1; raising μ (talent weight) raises the ceiling continuously within either regime. The largest mobility gains require both instruments applied jointly.

3.3 Small and Large Competitions: Where Policy Leverage Differs

The model's comparative statics on the competition size Q and the talent weight μ generate a sharp distinction between two empirically recognisable types of competition with different policy implications.

Small competitions (low Q , low μ). Partner-track interviews at professional services firms, senior academic appointments, and board-level selections involve a narrow field of challengers and assessors who can observe social capital directly: institutional affiliation, cultural fit, communication style, and network membership are all legible at close range. In the model's terms, these competitions have low Q and low μ (talent weight).

Two implications follow for this regime. First, the participation window A is relatively wide, and both Exclusive Participation (EP) and Full Participation (FP) are plausible depending on the level of $\bar{v}$. Cost-reduction policies have genuine traction here: since Q is small, the expected return to a low-income participant's investment is non-negligible, and reducing $\bar{v}$ can shift the equilibrium to universal participation. Second, this is the only regime in which the Inclusive Participation (IP) equilibrium — in which low-income participants over-invest while the high-income participant free-rides on raw talent — is empirically relevant. The probability surplus $\Phi(Q)$ is approximately 0.46 at $Q = 1$ and $\mu = 0.2$, giving a genuine feasibility window for the anomalous equilibrium. This is consistent with the documented tendency for aspiring candidates from lower-income backgrounds to over-invest in rank-building credentials — target university choices, internship accumulation, professional society memberships — in precisely the narrow-field selection processes where such credentials are most legible to assessors [Rivera, 2015].

The Inclusive Participation (IP) trap exists because μ (talent weight) is low: social capital is so decisive that low-income participants rationally over-invest in it even when the high-income participant abstains. Raising μ (talent weight) — making raw talent more decisive relative to institutional affiliation — directly collapses the trap. Since $\Phi(Q)$ is strictly decreasing in μ (talent weight), the feasibility window $U_1 < \Phi(Q) \leq U_2/Q$ narrows as μ (talent weight) rises, and the anomalous equilibrium disappears. Policy targeting small competitions should therefore combine cost reduction with active raising of μ (talent weight): introducing structured competency-based assessment criteria, reducing the weight placed on network affiliation in hiring decisions, and increasing transparency about what assessors actually value. Both instruments are necessary. Cost reduction alone — without raising μ (talent weight) — shifts low-income participants into a competition whose reward structure still favours social capital, potentially extending the over-investment trap to a larger pool rather than eliminating it.

Large competitions (high Q , higher μ). National civil service examinations, mass

graduate recruitment schemes, and standardised professional qualifications involve wide fields and assessors who must rely on criteria that are verifiable at scale. Standardisation is typically designed to increase μ (talent weight) — to make raw talent more decisive relative to institutional affiliation.

In this regime, two simultaneous forces operate. First, $\partial A / \partial Q < 0$ means the participation window narrows as the field grows: the range of costs sustaining universal participation shrinks regardless of policy intent, because each individual low-income participant’s expected return falls with Q . The direction of the equilibrium pressure is always toward elite-only participation as competition size increases.

Second, and working against this, raising μ (talent weight) through standardisation directly raises the ceiling $\mu / [2(1 - \mu)]$. Table 2 shows the magnitude: moving from $\mu = 0.20$ to $\mu = 0.40$ raises the ceiling from 12.5% to 33.3%. For large competitions where the equilibrium is likely to be Exclusive Participation (EP) regardless of cost levels, this is the operative lever. Access expansion without competition redesign reinforces elite-only equilibria: a larger pool drives A downward without offsetting the structural advantage of the high-income participant. Raising μ (talent weight) through standardisation and anonymisation is the primary instrument in this regime.

3.4 Implications for Education Policy

The model’s primary empirical application is to educational credentials and the competitions they determine entry into. Social capital in the model — the skills, credentials, and network affiliations unlocked by the threshold expenditure $\bar{v}$ — functions as a *productive input into tournament performance*: it directly raises the composite score $\mu r_i + (1 - \mu)s_i$ and thereby directly raises the probability of winning. This is a contest-theoretic mechanism in the tradition of Moldovanu et al. [2007], not a signaling mechanism in the tradition of Spence: there is no asymmetric information about pre-existing types, no separating equilibrium, and no off-path beliefs to be refined. The investment in social capital changes outcomes because it changes performance, not because it reveals anything about the investor. Two implications for education policy follow from this distinction.

Access policy and rank-building investment. In contest models, expanding access to the productive input raises competitive intensity without necessarily raising class-level mobility. In the present model, subsidising $\bar{v}$ — through means-tested bursaries, state-funded professional qualifications, or open-access credential programmes — expands the set of low-income participants who can afford to invest. If the subsidy is large enough to shift the equilibrium from EP to FP, the ceiling on collective mobility jumps from $\mu / [2(1 - \mu)]$ to the FP limit of 1: a discrete and substantial gain. If the subsidy is insufficient to shift the equilibrium — if it reduces $\bar{v}$ without moving it below y_1 — it raises competitive intensity within the EP regime without improving mobility outcomes. The distinction between margin-shifting and equilibrium-shifting subsidies is therefore crucial for policy evaluation, and it is invisible in models that treat participation as a continuous choice rather than a threshold decision. Because social capital is productive rather than merely informative, the case for credential subsidies rests on their equilibrium-shifting effect rather than on signaling arguments.

Competition design and the weight on social capital. The weight μ (talent weight) is a design parameter of the competition, not a fixed feature of the environment. In educational settings it corresponds to the relative weight placed on raw ability versus acquired credentials, institutional affiliation, and network membership in selection decisions. Standardised examinations, anonymised applications, and competency-based hiring frameworks all correspond formally to an increase in μ (talent weight): they reduce the contest advantage conferred by the threshold investment and make raw talent more decisive.

The mobility ceiling $\mu/[2(1 - \mu)]$ is directly determined by this design parameter. Because social capital is a productive input and not a signal, the policy instrument for raising the ceiling is competition redesign — changing what the contest rewards — rather than disclosure or information policy. This also resolves the tension between equality of opportunity and equality of outcome that Heckman and Mosso [2014] identify as central to education policy: the participation threshold $\bar{v}^*$ governs equality of opportunity and responds to $\bar{v}$ (the cost side); the mobility ceiling governs equality of outcome and responds to μ (talent weight, the design side). These are governed by separate parameters and require separately targeted instruments. Conflating them — treating credential subsidies as if they also raise the ceiling, or treating competition redesign as if it also removes budget constraints — risks improving one dimension without improving the other.

4 Conclusion

This paper develops a formal model of positional competition for elite credentials in the credentialist tradition: social capital is a rank-building input into status tournaments, not a human capital asset that accumulates across jobs. The barrier to entry is budgetary and threshold-structured. Wider income gaps simultaneously sustain high-income investment and price out low-income challengers, generating the empirical pattern — status expenditure rising with inequality — as the generic equilibrium outcome. The post-war prediction that narrowing income gaps would intensify rivalry is not wrong; it corresponds to a different equilibrium, one requiring cost levels low enough for universal participation.

Results. Four principal results emerge. First, under diminishing wealth utility, only two Nash equilibria are generically stable: universal participation (FP) and elite-only participation (EP). The transition is governed by a closed-form participation threshold $\bar{v}^*$ decreasing in both competition size Q and the social capital weight $(1 - \mu)$, with both forces pushing toward elite-only equilibria as competitions grow larger and credentials more decisive.

Second, collective low-income mobility converges to a hard ceiling of $\mu/[2(1 - \mu)]$ in the large- Q regime. This ceiling is bounded strictly below one-half, tightens as social capital is weighted more heavily, and is the operative social limit to growth as Hirsch (1977) originally conjectured.

Third, a reference-dependent equilibrium (IP) exists in which only low-income participants invest; it is viable only for $Q \leq 2$ and collapses rapidly as competition

size grows. This accounts for the documented over-investment in elite credentials by aspiring candidates from lower-income backgrounds in narrow selection processes, a phenomenon specific to small competitions and not replicable in large standardised assessments.

Fourth, competition size and competition design interact in a specific way: as Q grows, the equilibrium shifts toward elite-only participation at exactly the rate collective mobility converges to the ceiling. The two forces are not independent; they arrive together.

Empirical corroboration. Several predictions are already consistent with existing evidence. The Q -comparative static — that access expansion without cost reduction reinforces elite-only equilibria — is aligned with the Sutton Trust’s finding that the Russell Group share of disadvantaged students declined despite 25 years of widening participation policy [Sutton Trust, 2023]. It aligns also with China’s higher education expansion, which left access stratification essentially unchanged [Ding et al., 2021]. The mobility ceiling is consistent with Russell Group access rates of 1–1.4% against a theoretical equality benchmark of 4.4% — a directional and structural correspondence, not a direct numerical mapping. In another instance, Jerrim and Macmillan [2015] decompose the Great Gatsby Curve using PIAAC data to find that educational inequality accounts for a substantial share of the cross-national mobility gap, motivating access-widening as the primary policy instrument. The present model qualifies this conclusion in a specific way: the participation threshold $\bar{v}^*$ identifies the access channel they recover, but the model shows that reducing $\bar{v}^*$ alone — without simultaneously raising the talent weight μ — cannot move the mobility ceiling. The three empirical residuals introduced in the introduction — the dynastic wage penalty [Lo Bello and Morchio, 2022], the cross-occupational class pay gap [Laurison and Friedman, 2016], and the Italian liberalisation result [Mocetti et al., 2022] — each map directly onto the model’s μ (talent weight)– $(1 - \mu)$ (social capital weight) architecture, as set out there.

Policy implications. Two policy instruments emerge and their complementarity is the paper’s central applied finding. Reducing participation costs $\bar{v}$ — through credential subsidies, means-tested bursaries, open-access qualifications, or reduced network barriers — shifts the equilibrium toward universal participation when the reduction is large enough to move the budget constraint from binding to non-binding. This produces a discrete jump in the operative mobility ceiling from $\mu/[2(1 - \mu)]$ to 1. Raising the meritocratic weight μ through standardised assessment, anonymised selection, and competency-based criteria raises the ceiling directly. Cost reduction and competition redesign are complements: the former determines which equilibrium prevails, the latter determines the ceiling that bounds collective mobility under any equilibrium. These instruments are formally independent: the ceiling formula $\mu/[2(1 - \mu)]$ contains only μ (talent weight), so no reduction in $\bar{v}$, however large, can raise the ceiling without a simultaneous change in competition design. Moreover, sub-threshold reductions in $\bar{v}$ — those insufficient to move the budget constraint from binding to non-binding — raise Q without shifting the equilibrium, which reduces each individual low-income participant’s expected return and can tighten elite-only participation rather than loosen it.

This complementarity suggests why the long-term widening participation policy in

the United Kingdom — through Office for Students access targets, contextual admissions, and bursary programmes — has produced only modest improvements in elite mobility rates. These policies have primarily addressed $\bar{v}$ at the margin, producing small reductions in financial barriers without the equilibrium-shifting reductions required to move from EP to FP. Simultaneously, the design of elite competition — the weight placed on school type, extracurricular networks, and institutional affiliation in admissions and hiring — has remained largely unchanged, leaving μ (talent weight) and therefore the ceiling intact. The model implies that meaningful improvement in class-level mobility at elite institutions requires policy that addresses both parameters jointly, not sequentially.

The policy prescriptions differ between small and large competitions in the leverage each instrument provides, not in the direction of the μ (talent weight) prescription. In small competitions — partner-track interviews, senior academic appointments, narrow graduate schemes — both instruments have traction. Cost reduction shifts the equilibrium toward universal participation where Q is small enough to make individual expected returns non-negligible. And because the IP over-investment trap is driven by low μ (talent weight), raising μ (talent weight) directly collapses it: $\Phi(Q)$ is decreasing in μ (talent weight), so the feasibility window for the anomalous equilibrium narrows as talent becomes more decisive. Cost reduction alone — without also raising μ (talent weight) — risks extending the trap to a larger pool of entrants rather than eliminating it. The appropriate intervention combines both instruments: structured competency-based criteria, reduced weight on network affiliation, and transparent assessment rubrics jointly raise μ (talent weight); credential subsidies and open-access programmes reduce $\bar{v}$. In large competitions — civil service examinations, mass graduate recruitment, standardised professional qualifications — the participation window closes regardless of cost levels as Q grows. The operative lever is raising μ (talent weight) through standardisation and competition redesign. Expanding the applicant pool without redesigning what the competition rewards reinforces elite-only equilibria in this regime.

Because social capital is a productive rank-building input rather than a signal, the case for credential subsidies rests on their equilibrium-shifting effect, not entirely on signaling arguments. Credential subsidy programmes resisted on the grounds of signaling waste — the argument that expanded credential investment merely inflates the credential floor without generating social returns — might not always be the right theoretical framework to positional competitions where the investment directly raises competitive rank — situations that are even more likely to arise in growing economies. In positional contests, the welfare case for subsidies is that they shift the equilibrium to universal participation; it does not depend on whether credentials also signal productivity.

Future empirical work. The model generates testable predictions that go beyond what the existing literature has examined directly. The mobility ceiling $\mu/[2(1-\mu)]$ varies with μ (talent weight) but not with Q : cross-competition variation in the ceiling should be explained by observable proxies for credential weight in selection — the proportion of assessment based on institutional affiliation versus standardised tests, the weight placed on extracurricular social capital in hiring rubrics, the degree to which selection criteria are publicly observable and standardised — but should be orthogonal to competition size once credential weight is controlled for. This is an estimating equation

that can be taken to data using existing cross-institutional variation in admission and hiring criteria, and it distinguishes the present model from human capital models in which the ceiling depends on the productivity return to education rather than the contest design parameter.

The central message of the paper is that mobility cannot be expanded by widening access alone. The design of what credential competitions reward and who can afford to enter them are structurally distinct problems governed by distinct parameters — μ (talent weight) and $\bar{v}$ respectively — and the model provides the formal basis for treating them as separate policy objectives requiring separately targeted instruments. Conflating them, as twenty years of widening participation policy largely has, is not merely an oversight; it is a category error that the present framework makes precise.

References

- David B. Bills. Credentials, signals, and screens: Explaining the relationship between schooling and job assignment. *Review of Educational Research*, 73(4):441–449, 2003.
- Jo Blanden, Paul Gregg, and Stephen Machin. Intergenerational mobility in Europe and North America. Report supported by the Sutton trust, Centre for Economic Performance, London School of Economics, 2005. URL <https://www.suttontrust.com/wp-content/uploads/2020/01/IntergenerationalMobility.pdf>.
- Pierre Bourdieu. The forms of capital. In John G. Richardson, editor, *Handbook of Theory and Research for the Sociology of Education*, pages 241–258. Greenwood Press, New York, 1986.
- Pierre Bourdieu and Jean-Claude Passeron. *Reproduction in Education, Society and Culture*. Sage, London, 1977. Translated by Richard Nice.
- Andreas Chai, Wolfhard Kaus, and Christian Kiedaisch. Conspicuous spending and the income distribution of social groups. *Economic Inquiry*, 57(3):1324–1341, 2019.
- Harold L. Cole, George J. Mailath, and Andrew Postlewaite. Social norms, savings behavior, and growth. *Journal of Political Economy*, 100(6):1092–1125, 1992.
- Yanqing Ding, Yinduo Wu, Jin Yang, and Xiaoyang Ye. The elite exclusion: stratified access and production during the Chinese higher education expansion. *Higher Education*, 82(2):323–347, 2021. doi: 10.1007/s10734-021-00682-y.
- Robert H. Frank. Are workers paid their marginal products? *American Economic Review*, 74(4):549–571, 1984.
- Robert H. Frank. *Choosing the Right Pond: Human Behavior and the Quest for Status*. Oxford University Press, New York, 1985.
- John Kenneth Galbraith. *The Affluent Society*. Houghton Mifflin, Boston, 1958.

- Andrea Gallice and Edoardo Grillo. A model of educational investment, social concerns, and inequality. *The Scandinavian Journal of Economics*, 121(4):1620–1646, 2019. doi: 10.1111/sjoe.12299.
- James J. Heckman and Stefano Mosso. The economics of human development and social mobility. *Annual Review of Economics*, 6(1):689–733, 2014.
- Fred Hirsch. *Social Limits to Growth*. Routledge and Kegan Paul, London, 1977.
- Saravana Jaikumar and Ankur Sarin. Conspicuous consumption and income inequality in an emerging economy: evidence from India. *Marketing Letters*, 26(3):279–292, 2015. doi: 10.1007/s11002-015-9350-5.
- Christopher Jencks. *Who Gets Ahead? The Determinants of Economic Success in America*. Basic Books, New York, 1979.
- John Jerrim and Lindsey Macmillan. Income inequality, intergenerational mobility, and the Great Gatsby Curve: Is education the key? *Social Forces*, 94(2):505–534, 2015. doi: 10.1093/sf/sov075.
- Ye Jin, Hongbin Li, and Binzhen Wu. Income inequality, consumption, and social-status seeking. *Journal of Comparative Economics*, 39(2):191–204, 2011.
- Steven Jones. The personal statement: A fair way to assess university applicants? Sutton Trust Research Report, 2013. URL <https://www.suttontrust.com/wp-content/uploads/2019/12/JONESPERSONALSTATEMENTS-2.pdf>.
- Daniel Laurison and Sam Friedman. The class pay gap in Britain’s higher professional and managerial occupations. *American Sociological Review*, 81(4):668–695, 2016. doi: 10.1177/0003122416653602.
- Salvatore Lo Bello and Iacopo Morchio. Like father, like son: Occupational choice, intergenerational persistence and misallocation. *Quantitative Economics*, 13(2): 629–679, 2022. doi: 10.3982/QE1375.
- Jonathan J. B. Mijs. The paradox of inequality: income inequality and belief in meritocracy go hand in hand. *Socio-Economic Review*, 19(1):7–35, 2021.
- Jonathan J. B. Mijs and Mike Savage. Meritocracy, elitism and inequality. *Political Quarterly*, 91(2):397–404, 2020.
- Sauro Mocetti, Giacomo Roma, and Enrico Rubolino. Knocking on parents’ doors: Regulation and intergenerational mobility. *Journal of Human Resources*, 57(2): 525–554, 2022. doi: 10.3368/jhr.57.2.0219-10074R2.
- Benny Moldovanu, Aner Sela, and Xianwen Shi. Contests for status. *Journal of Political Economy*, 115(2):338–363, 2007.
- Andrew Postlewaite. The social basis of interdependent preferences. *European Economic Review*, 42(3–5):779–800, 1998.

Lauren A. Rivera. *Pedigree: How Elite Students Get Elite Jobs*. Princeton University Press, Princeton, 2015.

Sutton Trust. Little progress in opening up top universities to poorer pupils in 25 years. Sutton Trust Report, October 2023. URL <https://www.suttontrust.com/news-opinion/all-news-opinion/little-progress-in-opening-up-top-universities-to-poorer-pupils>.

Bernd Wegener. Concepts and measurement of prestige. *Annual Review of Sociology*, 18:253–280, 1992. doi: 10.1146/annurev.so.18.080192.001345.

A Convolution of Uniform Random Variables

The distribution of a convex combination $w = \mu r + (1 - \mu)s$ of two random variables r, s - which are distributed according to densities $f_R(r)$ and $f_S(s)$ - can be derived using the convolution of two derived variables $x = \mu r$ and $y = (1 - \mu)s$. If r and s are uniformly distributed variables ranging from 0 to 1, then x ranges from 0 to μ and y ranges from 0 to $(1 - \mu)$. The CDF for $F_X(x)$ and $F_Y(y)$ are simply $F_X(x) = \frac{x}{\mu}$, $F_Y(y) = \frac{y}{1-\mu}$ (respectively). Therefore, the densities for x and y are $f_X(x) = \frac{1}{\mu}$ and $f_Y(y) = \frac{1}{1-\mu}$. The density of w is simple $f(w) = \int_{-\infty}^{\infty} f_Y(w-x)f_X(x)dx$. Since, $f_X(x)$ is $1/\mu$ only between 0 and μ . We have the following integral to evaluate for the convolution

$$f(w) = \frac{1}{\mu} \int_0^{\mu} f_Y(w-x)dx$$

Since $f_Y(y)$ is 1 only between 0 and $1 - \mu$, we have the condition $0 < w - x < 1 - \mu \Rightarrow w - (1 - \mu) < x < w$. Further, x is distributed only between 0 and μ . As shown in Figure 2, the integral would evaluate to the following⁵ under the assumption $\mu < \frac{1}{2}$.

$$f(w) = \frac{1}{\mu} \int_0^{\mu} f_Y(w-x)dx = \begin{cases} \frac{1}{\mu} \int_0^w \frac{1}{1-\mu} dx & 0 < w < \mu \\ \frac{1}{\mu} \int_0^{\mu} \frac{1}{1-\mu} dx & \mu < w < 1 - \mu \\ \frac{1}{\mu} \int_{w-(1-\mu)}^{\mu} \frac{1}{1-\mu} dx & 1 - \mu < w < 1 \\ 0 & \text{otherwise} \end{cases}$$

This can be simplified as

$$f(w) = \begin{cases} \frac{w}{\mu(1-\mu)} & 0 < w < \mu \\ \frac{1}{1-\mu} & \mu < w < 1 - \mu \\ \frac{1-w}{\mu(1-\mu)} & 1 - \mu < w < 1 \\ 0 & \text{otherwise} \end{cases}$$

⁵Notice that as a convex combination, w is distributed only between 0 and 1.

Figure 2 shows this trapezoidal density with μ set to $\frac{1}{5}$. The CDF $F_X(x)$ can be evaluated by integrating the above density as $P(X < x) = \int_{-\infty}^x f(t)dt$.

$$F_X(x) = \begin{cases} \frac{x^2}{2\mu(1-\mu)} & 0 < x < \mu \\ \frac{\mu}{2(1-\mu)} + \frac{x-\mu}{1-\mu} & \mu < x < 1-\mu \\ \frac{\mu^2}{2\mu(1-\mu)} + \frac{2\mu(1-2\mu)}{2\mu(1-\mu)} + \frac{\mu^2-x^2+2x-1}{2\mu(1-\mu)} & 1-\mu < x < 1 \\ 0 & \text{otherwise} \end{cases}$$

Simplifying above, we have

$$F_X(x) = \begin{cases} \frac{x^2}{2\mu(1-\mu)} & 0 < x < \mu \\ \frac{x-\mu/2}{1-\mu} & \mu < x < 1-\mu \\ 1 - \frac{(x-1)^2}{2\mu(1-\mu)} & 1-\mu < x < 1 \\ 0 & \text{otherwise} \end{cases}$$




Figure 2: Left: ranges of w and x for convolution of two scaled random variables. Right: trapezoidal density of w as the convex combination with $\mu = \frac{1}{5}$.

B Winning Probability Derivations

EP: High-income invests, low-income do not

The low-income participant's score is μr_i (no social capital). She wins if μr_i exceeds all $Q - 1$ other low-income participants' scores and the high-income participant's full composite score $x_H = \mu r_H + (1 - \mu)s_H$. Since $\mu r_i < \mu$ for all $r_i \in (0, 1)$ under Assumption 1, only the first piece of F_X applies:

$$P_1(\text{win}|r_i) = r_i^{Q-1} \cdot F_X(\mu r_i) = r_i^{Q-1} \cdot \frac{\mu r_i^2}{2(1-\mu)} = \frac{\mu}{2(1-\mu)} r_i^{Q+1}$$

Integrating over $r_i \sim \text{Uniform}[0, 1]$:

$$p_1 = \int_0^1 \frac{\mu}{2(1-\mu)} r^{Q+1} dr = \frac{\mu}{2(1-\mu)(Q+2)}$$

IP: Low-income invest, high-income does not

A given low-income participant wins if her full score $x_i = \mu r_i + (1 - \mu)s_i$ exceeds x_k for all $Q - 1$ other low-income participants and exceeds μr_H for the non-investing high-income participant. The conditional probability is:

$$P_1(\text{win}|r_i, s_i) = F_X(x_i)^{Q-1} \cdot G(x_i), \quad G(x) = \begin{cases} x/\mu & 0 < x < \mu \\ 1 & x \geq \mu \end{cases}$$

For $Q = 1$ the double integral over the unit square yields:

$$p'_1 = \frac{6-7\mu}{6(1-\mu)} \quad (Q = 1)$$

For $Q \geq 2$, $F_X(x)^{Q-1} \cdot G(x)$ produces a rational function of μ of degree Q within each CDF region; each specific Q yields a closed form but no compact general formula holds uniformly across all Q .

FW and FP

When neither invests, scores are μr_i ; by symmetry each of $Q + 1$ participants wins with probability $1/(Q + 1)$:

$$p_1 = \int_0^1 r^Q dr = \frac{1}{Q+1}$$

When both invest, all scores x_i are identically distributed via F_X ; by symmetry $p_1 = 1/(Q + 1)$ again.

Integration regions

To calculate $\iint P(\text{win}|r,s) dr ds$ over the unit square, we split by the boundaries of the piecewise function. The three regions are shown in Figure 3. If $x = \mu r + (1 - \mu)s$ and $F(x)$ is of the piecewise form with pieces A, B, C :

$$\iint F(x) ds dr = \int_{r=0}^1 \int_0^{\frac{\mu(1-r)}{1-\mu}} A ds dr + \int_{r=0}^1 \int_{\frac{\mu(1-r)}{1-\mu}}^{1-\frac{\mu r}{1-\mu}} B ds dr + \int_{r=0}^1 \int_{1-\frac{\mu r}{1-\mu}}^1 C ds dr$$




Figure 3: Integration over r, s split by boundaries $\mu r + (1 - \mu)s = \mu$ and $\mu r + (1 - \mu)s = 1 - \mu$

Asymptotic probability surplus

Proposition 2 (Asymptotic probability surplus).

$$\lim_{Q \rightarrow \infty} Q(Q+1)\Phi(Q) = 1$$

equivalently, $\Phi(Q) = \frac{1}{Q(Q+1)} [1 + O(1/Q^2)]$. The convergence is $O(1/Q^3)$ and independent of μ .

Proof. Write $\Phi(Q) = \int_0^1 F_X(t)^{Q-1} [G(t) - F_X(t)] f_X(t) dt$, where the identity follows from $p'_1 = \int F_X^{Q-1} G f_X dt$ and $1/(Q+1) = \int F_X^Q f_X dt$. As $Q \rightarrow \infty$, $F_X(t)^{Q-1}$ concentrates near $t = 1$. Setting $u = 1 - t$ and $c = \mu(1 - \mu)$: $F_X \sim 1 - u^2/(2c)$, $G - F_X \sim u^2/(2c)$, $f_X \sim u/c$. By Laplace's method:

$$\Phi(Q) \sim \frac{1}{2c^2} \int_0^\infty u^3 \exp\left(-\frac{Qu^2}{2c}\right) du = \frac{1}{2c^2} \cdot \frac{2c^2}{Q^2} = \frac{1}{Q^2}$$

Since $c = \mu(1 - \mu)$ cancels, the leading term is μ -independent. The correction is $O(1/Q^3)$. $\square$

Corollary 1. $Q \cdot \Phi(Q) \rightarrow 0$ at rate $1/Q$. The IP feasibility condition $U_1 < \Phi(Q)$ fails for all $U_1 > 0$ for sufficiently large Q : IP is asymptotically impossible.

Table 3 reports $\Phi(Q)$ for $Q = 1, \dots, 10$ at representative values of μ , confirming the rapid convergence and μ -independence.

Table 3: Probability surplus $\Phi(Q)$ and asymptotic approximation $1/[Q(Q+1)]$

Q	$\mu = 0.1$	$\mu = 0.2$	$\mu = 0.3$	$\mu = 0.4$	$1/[Q(Q+1)]$	$Q(Q+1)\Phi$
1	0.4815	0.4583	0.4286	0.3889	0.5000	0.917
2	0.1664	0.1651	0.1621	0.1556	0.1667	0.991
3	0.0833	0.0832	0.0829	0.0816	0.0833	0.998
4	0.0500	0.0500	0.0499	0.0497	0.0500	1.000
5	0.0333	0.0333	0.0333	0.0333	0.0333	1.000
6	0.0238	0.0238	0.0238	0.0238	0.0238	1.000
7	0.0179	0.0179	0.0179	0.0179	0.0179	1.000
8	0.0139	0.0139	0.0139	0.0139	0.0139	1.000
9	0.0111	0.0111	0.0111	0.0111	0.0111	1.000
10	0.0091	0.0091	0.0091	0.0091	0.0091	1.000

Values computed via numerical double integration. The last column reports $Q(Q+1)\Phi(Q)$ at $\mu = 0.2$; values at other μ are identical to the precision shown for $Q \geq 4$.

C T-Period Utility Generalisation

The main text presents intertemporal utility under T periods with discount sum $D_T = \sum_{t=0}^{T-2} \delta^{-t}$. We establish here that the qualitative equilibrium structure is unchanged by T .

Closed form for D_T . Since $\delta > 1$, the geometric sum evaluates to:

$$D_T = \frac{1 - \delta^{-(T-1)}}{1 - \delta^{-1}}$$

For $T = 2$: $D_T = 1$, recovering the two-period model. For $T = 3$: $D_T = 1 + 1/\delta$.

Proposition (T-period invariance). *The set of Nash equilibria under T -period utility is identical to that under two-period utility ($T = 2$). The participation threshold $\bar{v}^*$ scales monotonically with D_T : it is strictly increasing in T and strictly decreasing in δ .*

Proof. All Nash equilibrium conditions (Equations 8–13) share the structure:

$$\frac{D_T}{\delta} \cdot \Delta p \cdot (u(y_2) - u(y_1)) \geq U_i$$

where Δp is a probability gap and U_i is a utility cost. Since D_T/δ is a common positive scaling factor, it does not alter which side of any inequality is larger: an equilibrium condition satisfied under $T = 2$ is satisfied for all $T \geq 2$ with the same qualitative direction. The threshold $\bar{v}^*$ (Equation 23) has $A \propto D_T/\delta$; since $(y_1/y_2)^A$ is decreasing in A and $y_1 < y_2$, $\bar{v}^* = y_1(1 - (y_1/y_2)^A)$ increases with A , hence with D_T . D_T is strictly increasing in T (adding positive terms) and strictly decreasing in δ (each term δ^{-t} is decreasing in δ). $\square$

The proposition implies that a longer investment horizon or greater patience raises $\bar{v}^*$, expanding the range of costs under which universal participation (FP) is sustained. The two-period model is therefore a conservative baseline: all comparative statics on $\bar{v}^*$ with respect to μ and Q carry over unchanged to the T -period case.

Log utility closed form

Under log utility $u(y) = \log(y)$, the participation threshold $\bar{v}^*$ has a tractable closed form. The indifference condition $U_1 = \frac{D_T}{\delta} \left(\frac{1}{Q+1} - p_1 \right) (u(y_2) - u(y_1))$ becomes:

$$(22) \quad \log \frac{y_1}{y_1 - \bar{v}^*} = A \log \frac{y_2}{y_1}$$

where $A = \frac{D_T}{\delta} \left(\frac{1}{Q+1} - \frac{\mu}{2(1-\mu)(Q+2)} \right) > 0$. Solving:

$$(23) \quad \bar{v}^* = y_1 \left(1 - \left(\frac{y_1}{y_2} \right)^A \right)$$

Since $(y_1/y_2)^A$ is decreasing in A and $y_1 < y_2$, $\bar{v}^*$ is increasing in A , hence decreasing in μ and Q (since $\partial A/\partial \mu < 0$ and $\partial A/\partial Q < 0$). The threshold depends on the income ratio y_1/y_2 rather than the income level alone.

D Utility Payoffs

Table 4 reports the full utility payoffs for each of the four investment cases. In each cell, the upper entry is the high-income participant's utility and the lower entry is the low-income participant's utility. The scale factor D_T/δ multiplies all second-period terms.

Table 4: Utility payoffs (scale factor D_T/δ multiplies all second-period terms)

		<i>Low-income withdraws</i>	<i>Low-income invests</i>
<i>High-income withdraws</i>	High-income	$u(y_2) + \frac{D_T}{\delta} \left[\frac{u(y_2)}{Q+1} + \frac{Qu(y_1)}{Q+1} \right]$	$u(y_2) + \frac{D_T}{\delta} [(1 - Qp'_1)u(y_2) + Qp'_1u(y_1)]$
	Low-income	$u(y_1) + \frac{D_T}{\delta} \left[\frac{u(y_2)}{Q+1} + \frac{Qu(y_1)}{Q+1} \right]$	$u(y_1 - \bar{v}) + \frac{D_T}{\delta} [p'_1u(y_2) + (1 - p'_1)u(y_1)]$
<i>High-income invests</i>	High-income	$u(y_2 - \bar{v}) + \frac{D_T}{\delta} [p_2u(y_2) + (1 - p_2)u(y_1)]$	$u(y_2 - \bar{v}) + \frac{D_T}{\delta} \left[\frac{u(y_2)}{Q+1} + \frac{Qu(y_1)}{Q+1} \right]$
	Low-income	$u(y_1) + \frac{D_T}{\delta} [p_1u(y_2) + (1 - p_1)u(y_1)]$	$u(y_1 - \bar{v}) + \frac{D_T}{\delta} \left[\frac{u(y_2)}{Q+1} + \frac{Qu(y_1)}{Q+1} \right]$

p_1, p_2 : EP probabilities. p'_1 : IP low-income probability ($Q = 1$ closed form; general Q numerical). The high-income participant's IP probability satisfies $1 - Qp'_1$, using the identity $1/(Q+1) - (1 - Qp'_1) = Q(p'_1 - 1/(Q+1))$, valid for all $Q \geq 1$.

Corresponding author: Anurag Srivastava, Riskcare Ltd., anurag.srivastava@riskcare.com.
ORCID: 0000-0002-6477-4430.