

Sequence effects at independence under partial adoption

Setting. The prior covariance is $c = 0$, so the two attributes are believed unrelated, and the adoption weight satisfies $0 < \omega < 1$, so the cue read second is adopted in part. At $c = 0$ the paper's channel, sequence dependence carried through the prior association, is switched off. What remains is a second channel, the weakening of whichever cue is read second, which is at full strength here. Full adoption ($\omega = 1$) is the setting of Sections 4 and 5, where the second channel is absent and the first vanishes at $c = 0$ (Proposition IMM). All quantities below are exact at $c = 0$; the $\mathcal{O}(c)$ corrections are in the script and are not the point.

Table 1: What can be said at $c = 0$ and $0 < \omega < 1$.

	Statistic	Value at $c = 0$	What it says
<i>Between the two sequences, one evaluator</i>			
1	marginal probabilities	$P_{AB}^J(A=1) - P_{BA}^J(A=1) = (1 - \omega)(\alpha - q_0)$ $P_{AB}^J(B=1) - P_{BA}^J(B=1) = -(1 - \omega)(\beta - r_0)$	The marginals carry the whole effect. Each is the distance the cue would have moved its marginal, scaled by how much of it was not adopted. Zero only at $\omega = 1$ or when the cue delivers the prior marginal. The mechanism is position, not association: in AB the letter is the weakened cue, in BA the credential is.
2	believed association (cross-product; odds ratio, Yule's Q)	$\text{assoc}(P_{AB}^J) - \text{assoc}(P_{BA}^J) = 0$ odds ratio identical, for every c	Both sequences leave a product table at $c = 0$, so the cross-product is zero in each; the odds ratio equals the prior's because every damped step is a row or column rescaling (Lemma SEP). The paper's qualitative verdict survives: the marginals see the sequence, the association does not. What changes is the order, zeroth instead of first, and with it the relevance of small c .
3	the adoption weight	$\omega = 1 - \frac{P_{AB}^J(A=1) - P_{BA}^J(A=1)}{\alpha - q_0}$	Row 1 read backwards. The channel the paper does not study is the channel Proposition ADJ reads; the identification device lives entirely at this point. Requires $q_0 \neq \alpha$.
4	decision flip and loss	score gap = $\langle v, P_{AB}^J - P_{BA}^J \rangle$, a fixed multiple of row 1, independent of c	The share of evaluators whose decision flips is the mass of a band of fixed width, not $\mathcal{O}(c)$, and the surplus-weighted loss is not second order. <i>Inferred from linearity of the score; not separately verified.</i>

Pooled over a population, against the benchmark, $0 < \lambda < 1$

	Statistic	Value at $c = 0$	What it says
5	association of the mean belief	$\text{assoc}(\overline{P}_\lambda) - \text{assoc}(P^B) = -\lambda(1-\lambda)(1-\omega)^2(\alpha - q_0)(\beta - r_0)$	The mean belief acquires an association nobody holds. Every evaluator's table is independent and so is the benchmark's, but the mixture of two product tables with different marginals is not a product table. A pooled audit of whether the traits go together returns a nonzero value in a population where no one believes they do. The sign is fixed by whether the two cues move their marginals in the same direction relative to the prior. No counterpart at $\omega = 1$, where the mixture at $c = 0$ is $q \otimes r$.
6	marginals of the mean belief	$\overline{P}_\lambda(A=1) - P^B(A=1) = (1-\lambda)(1-\omega)(q_0 - \alpha)$	Zeroth order. P^B sets both marginals to their delivered credences, so it is a full-adoption object under the other reading, and the gap from it mixes two differences, the reading and the adoption weight. No sequence-free reference exists at interior ω under either reading: a Bayes-factor update that scales the second cue's factor by ω is itself order-dependent at $c = 0$ (script, row 5), because Wagner's commutation needs the same factor in either position. The between-sequence comparison of rows 1–2 needs no reference, which is one more reason to prefer Proposition ORD's form of the test.

The two settings compared. The primacy regime is taken for the interior, since with ω on the second cue every interior value weakens the later impression. Weisberg (*Commutativity or Holism?*, read in full 2026-09-27; see `literature/weisberg2009`) supplies the vocabulary. He separates the *rule* that propagates a change on a partition from the *inputs problem* of how an experience fixes the values on that partition; Lange's point is that reversing two experiences does not reverse their input values; Wagner's theorem, read as a dilemma, says that commutativity on experiences forces each experience to carry the same Bayes factor whatever background it meets, which is anti-holism; and rigidity, the holding of conditionals fixed, prevents a later discovery from ever *undercutting* earlier support, leaving only rebutting. Marks: † a row of Table 1 or of `verify_interior_omega.py`; ‡ inferred from those rows, not separately checked; ◇ a mapping onto Weisberg's terms, not a computation.

Table 2: Full adoption against partial adoption (primacy regime).

	Full adoption, $\omega = 1$	Partial adoption, $0 < \omega < 1$
Where the order dependence arises ◇	In the rule. The inputs q, r are the same in either position; order matters because two rigid steps on correlated partitions do not commute (Diaconis–Zabell).	In the inputs. The value the second cue delivers depends on its position; the rule adds nothing new at $c = 0$.
Source in the model †	Association only. Zero at $c = 0$ (IMM).	Position, present at $c = 0$, plus association.

	Full adoption, $\omega = 1$	Partial adoption, $0 < \omega < 1$
Status of the inputs $\diamond$	Order-independent. Anti-holistic at the inputs, in Weisberg's sense: the credential delivers q whether or not the letter was read.	Order-dependent. History-dependent inputs, not holism in Weisberg's sense: the import of the second cue depends on its position, not on the credence in any defeater proposition.
Weisberg's dilemma, Bayes-factor reading $\diamond$	The commutative horn. Field's α , Wagner's theorem, the benchmark P^B .	Off the commutative horn, but not on the holistic one. A position-dependent factor is order-dependent at $c = 0$ (row 5 of the script) because commutation needs the same factor in either position; the dependence is on position, not on a defeater belief.
Weisberg's dilemma, delivered-credence reading $\diamond$	On neither horn. Inputs fixed regardless of any background, more anti-holistic than Field; non-commutative because the implied factor $r_j/P^A(B=j)$ varies with what the first cue did. This is non-commutativity on input distributions, the case Lange declines and Weisberg sets aside (p. 9). The two readings differ as much as they ever do.	Non-commutative through the inputs, as under the other reading; the readings differ only in how the surviving part of the earlier impression is blended, linearly in probability or geometrically in odds.
Rigidity $\diamond$	Both settings are rigid. Rigidity preserves independence, which is IMM at $c = 0$ for every ω ; a later cue can never undercut what an earlier one implied, only rebut it, and the only rebutting route is the prior association c .	
Size of the effect $\dagger$	Small, first order in c , spread unevenly across statistics.	Blunt. $(1 - \omega)(\alpha - q_0)$ on a marginal whatever c is.
Marginals $\dagger$	Last-read exact, first-read drifts at first order (DRF).	Both differ at zeroth order between sequences.
Believed association $\dagger$	Second order, between sequences and pooled (DEC, PRO).	Zero between sequences at $c = 0$, first order in c . Pooled, nonzero at $c = 0$, $-\lambda(1 - \lambda)(1 - \omega)^2(\alpha - q_0)(\beta - r_0)$.
Odds ratio $\dagger$	Identical across sequences, every c .	Identical across sequences, every c . The one statistic that does not move.
Decisions $\ddagger$	Share of flips first order, loss second order (SHR, LOS).	Score gap fixed in c , so the share is a fixed band and the loss is not second order.
Classification by order in c $\dagger$	Is the contribution. Which statistics are protected, and why.	Is moot. Marginal-based statistics see the effect at zeroth order, the odds ratio never; there is no fine structure to classify.
Audit implication $\dagger$	False negative. An association audit passes a sequence-dependent population.	False positive as well. A pooled association audit returns a stereotype no member holds, from mixing alone.
Identification	The reading, from which marginal ignores the prior association (ORD, DRF).	The weight, from three ratings (ADJ). At its strongest here, since the position channel is what ADJ reads.

	Full adoption, $\omega = 1$	Partial adoption, $0 < \omega < 1$
Literature stance $\diamond$	Diaconis–Zabell’s premise; Hawthorne’s objection inherited; Lange’s point declined.	Hawthorne answered by measurement; Hogarth–Einhorn’s anchoring is the model; Lange’s point conceded by one parameter; commutativity given up on principle, not only through the reading.

What the comparison says about scope. The paper’s contribution lives in the regime where the sequence effect is subtle, arising in the rule rather than in the inputs, existing only because the traits are believed related, and therefore small and spread unevenly across statistics. That is the regime in which an audit can be fooled by choosing the wrong statistic. In the primacy regime the effect is blunt and any marginal shows it; the question becomes not whether the sequence registers but how much of the earlier impression survived, which is a question about ω , answered by ADJ and nothing else. The two regimes are two different questions. Under full adoption the association audit fails by silence; under partial adoption it fails by fabrication as well.

One sentence for the paper. Full adoption is the setting in which the only source of sequence dependence is the one the paper studies, the cross-attribute effect carried through the prior association, which vanishes at $c = 0$; partial adoption adds a second source, the weakening of whichever cue is read second, which is present at $c = 0$ and so lies outside the expansion in c on which the results rest.